

Contents

| | | |
|----------|--|-----------|
| 1 | General Introduction | 2 |
| 2 | Short introduction to parallel Haskells | 2 |
| 2.1 | Multicore Haskell | 2 |
| 2.2 | ParMonad | 2 |
| 2.3 | Eden | 2 |
| 2.4 | HdpH | 2 |
| 3 | Arrows | 2 |
| 4 | Utility Functions | 3 |
| 5 | Parallel Arrows | 4 |
| 5.1 | Multicore Haskell | 5 |
| 5.2 | ParMonad | 5 |
| 5.3 | Eden | 6 |
| 5.4 | HdpH | 6 |
| 6 | Skeletons | 6 |
| 6.1 | parEvalNLazy | 6 |
| 6.2 | parEval2 | 7 |
| 6.3 | parMap | 7 |
| 6.4 | parMapStream | 8 |
| 6.5 | farm | 8 |
| 6.6 | farmChunk | 9 |
| 7 | Syntactic Sugar | 9 |
| 8 | Benchmarks | 10 |
| 9 | Conclusion | 10 |

1 General Introduction

Arrows were introduced in John Hughes paper as an alternative to Monads for API design.¹In the paper Hughes describes that Arrows have one powerful additional property when compared to Monads for API design: Extensibility. In this paper we will show how this property can be used to add parallelism capabilities to different arrows.

2 Short introduction to parallel Haskells

2.1 Multicore Haskell

2.2 ParMonad

2.3 Eden

2.4 HdpH

3 Arrows

John Hughes defined the Arrow typeclass as follows²:

```
1 class Arrow a where
2   arr :: (b -> c) -> a b c
3   (>>>) :: a b c -> a c d -> a b d
4   first :: a b c -> a (b,d) (c,d)
```

`arr :: (b -> c) -> a b c` is used to lift an ordinary function to an Arrow type. This can be thought of as analogous to the monadic **return**. The `>>>` operator, in a similar way, is analogous to the monadic composition operator `>>=`. Ant lastly, the `first` operator, which takes the input arrow from `b` to `c` and converts it into an arrow on pairs with the second argument untouched, is also needed for actual useful code as without it, we wouldn't have a way to save input across arrows.

The most prominent instances of this interface are regular functions, `(->)`

```
1 instance Arrow (->) where
2   arr f = f
3   f >>> g = g . f
4   first f = \ (b, d) -> (f b, d)
```

and the Kleisli type.

```
1 instance Arrow (Kleisli m) where
2   arr f = Kleisli $ return . f
```

¹John Hughes, Generalising Monads to Arrows, see [1]

²John Hughes, Generalising Monads to Arrows, see [1]

```

3 f >>> g = Kleisli $ \b -> f b >>= g
4 first f = Kleisli $ \ (b,d) -> f b >>= \c -> return (c,d)

```

With this typeclass in place, Hughes also defined some syntactic sugar: The mirrored version of `first`, called `second`,

```

1 second :: Arrow a => a b c -> a (d, b) (d, c)
2 second f = arr swap >>> first f >>> arr swap
3 where swap (x, y) = (y, x)

```

the `***` combinator which combines `first` and `second` to handle two inputs in one arrow,

```

1 (***) :: Arrow a => a b c -> a d e -> a (b, d) (c, e)
2 f *** g = first f >>> second g

```

and the `&&&` combinator that constructs an arrow which outputs 2 different values like `***`, but takes only one input.

```

1 (&&&) :: Arrow a => a b c -> a b d -> a b (c, d)
2 f &&& g = arr (\b -> (b, b)) >>> (f *** g)

```

A short example given by Hughes on how to use this is `add` over arrows:

```

1 add :: Arrow a => a b Int -> a b Int -> a b Int
2 add f g = (f &&& g) >>> arr \(u, v) -> u + v

```

The benefit of this interface is now that any type which is shown to be an `Arrow` can no be used in conjunction with this newly created `add` combinator. Even though this example is quite simple, the power of the `Arrow` interface immediately is clear: If a type is an `Arrow`, it can immediately used together with every library that works on `Arrows`. With simple `Monads` this type of extensibility is not possible.

Note: In the definitions above we used the notation `a b c` for an arrow from `b` to `c`. From now on we will use the equivalent definition `arr a b` for an arrow from `a` to `b` to make it easier to find the arrow type in type signatures. We kept the original notation `a b c` for this section to not change too much from Hughes' original definitions.

4 Utility Functions

Before we go into detail on parallel arrows, we introduce some utility combinators first, that will help us later: `map` and `zipWith` on arrows.

The `map` combinator lifts any arrow `arr a b` to an arrow `arr [a] [b]`,

```

1 -- from http://www.cse.chalmers.se/~rjmh/afp-arrows.pdf
2 mapArr :: ArrowChoice arr => arr a b -> arr [a] [b]
3 mapArr f =
4   arr listcase >>>

```

```

5 arr (const []) ||| (f *** mapArr f >>> arr (uncurry ()))
6 where
7   listcase [] = Left ()
8   listcase (x:xs) = Right (x,xs)

```

and zipWith lift any arrow `arr (a, b) c` to an arrow `arr ([a], [b]) [c]`.

```

1 zipWithArr :: ArrowChoice arr => arr (a, b) c -> arr ([a], [b]) [c]
2 zipWithArr f = (arr $ \ (as, bs) -> zipWith (,) as bs) >>> mapArr f

```

These two combinators make use of the `ArrowChoice` typeclass, which allows us to use the `|||` combinator, which takes two arrows `arr a c` and `arr b c` and combines them into a new arrow `arr (Either a b) c` which pipes all **Left** a's to the first arrow and all **Right** b's to the second arrow.

```

1 (|||) :: ArrowChoice arr a c -> arr b c -> arr (Either a b) c

```

With the `zipWithArr` combinator we can also write a combinator `listApp`, which lifts a list of arrows `[arr a b]` to an arrow `arr [a] [b]`.

```

1 listApp :: (ArrowChoice arr, ArrowApply arr) => [arr a b] -> arr [a] [b]
2 listApp fs = (arr $ \ as -> (fs, as)) >>> zipWithArr app

```

This combinator also makes use of the `ArrowApply` typeclass which allows us to evaluate arrows with `app :: arr (arr a b, a) c`.

Do we need the exact definition of `ArrowChoice` here?

5 Parallel Arrows

We have seen what Arrows are, so we can take a look at how they can be used to define a general interface not just to computation, but to **parallel computation** as well, next.

In its purest form, parallel computation (on functions) can be looked at as the execution of some functions `a -> b` in parallel:

```

1 parEvalN :: [a -> b] -> [a] -> [b]

```

Translating this into arrow terms gives us a new operator `parEvalN` that lifts a list of arrows `[arr a b]` to an parallel arrow `arr [a] [b]` (This combinator is similar to `listApp` from 4, which does no parallel evaluation of `[arr a b]`).

```

1 parEvalN :: [arr a b] -> arr [a] [b]

```

As the implementation may a) differ depending on the type of arrow and b) we want this to be an interface for different backends with, we introduce the new typeclass `ArrowParallel`.

```

1 class Arrow arr => ArrowParallel arr a b where
2   parEvalN :: [arr a b] -> arr [a] [b]

```

As computation in some parallel Haskell involves additional configuration parameters giving info about the environment, we also introduce an additional `conf` parameter to the function which can be of different type for each backend, so we add it to the type signature of our class as well.

```
1 class Arrow arr => ArrowParallel arr a b conf where
2   parEvalN :: conf -> [arr a b] -> arr [a] [b]
```

Note that we don't require the `conf` parameter in every implementation. If it is not needed, we allow the `conf` type parameter to be of any type and don't even evaluate it by blanking it in the type signature of the implemented `parEvalN`.

5.1 Multicore Haskell

The Multicore Haskell implementation of this class is straightforward using `listApp` from 4 combined with the `using` operator from Multicore Haskell.

```
1 instance (NFData b, ArrowApply arr, ArrowChoice arr) =>
2   ArrowParallel arr a b conf where
3   parEvalN _ fs = listApp fs >>> arr (flip using $ parList rdeepseq)
```

We hardcode the `parList rdeepseq` strategy here as in this context it is the only one making sense as we usually want the output list to be fully evaluated to its normal form.

5.2 ParMonad

The `ParMonad` implementation makes use of Haskell's laziness and `ParMonad`'s `spawnP :: a -> Par (IVar a)` function which forks away the computation of a value and returns an `IVar` containing the result in the `Par` monad.

We therefore apply each function to its corresponding input value with `app` and then fork the computation away with `arr spawnP` inside a `zipWithArr` call. This yields a list `[Par (IVar b)]`, which we then convert into `Par [IVar b]` with `arr sequenceA`. In order to wait for the computation to finish, we map over the `IVars` inside the `ParMonad` with `arr (>>= mapM get)`. The result of this operation is a `Par [b]` from which we can finally remove the monad again by running `arr runPar` to get our output of `[b]`.

```
1 instance (NFData b, ArrowApply arr, ArrowChoice arr) =>
2   ArrowParallel arr a b conf where
3   parEvalN _ fs =
4     (arr $ \as -> (fs, as)) >>>
5     zipWithArr (app >>> arr spawnP) >>>
6     arr sequenceA >>>
7     arr (>>= mapM get) >>>
8     arr runPar
```

5.3 Eden

For the Multicore and ParMonad implementation we could use general implementations that just required the ArrowApply and ArrowChoice interface. With Eden this is not the case as we can only spawn a list of functions and we cannot extract simple functions out of arrows. While this could be "fixed" by introducing a typeclass like

```
1 class (Arrow arr) => ArrowUnwrap arr where
2   arr a b -> (a -> b)
```

we don't do this in this paper, as this seems too hacky. For now, we just implement ArrowParallel for normal functions

```
1 instance (Trans a, Trans b) => ArrowParallel (->) a b conf where
2   parEvalN _ fs as = spawnF fs as
```

and the Kleisli type.

```
1 instance (Monad m, Trans a, Trans b, Trans (m b)) =>
2   ArrowParallel (Kleisli m) a b conf where
3   parEvalN conf fs =
4     (arr $ parEvalN conf (map (\(Kleisli f) -> f) fs)) >>>
5     (Kleisli $ sequence)
```

5.4 Hdph

6 Skeletons

With the ArrowParallel typeclass in place and implemented, we can now implement some basic parallel skeletons.

6.1 parEvalNLazy

parEvalN is 100% strict, which means that it fully evaluates all passed arrows. Sometimes this might not be feasible, as it will not work on infinite lists of functions like e.g. `map (arr . (+)) [1..]` or just because we need the arrows evaluated in chunks. `parEvalNLazy` fixes this. It works by first chunking the input from `[a]` to `[[a]]` with the given `ChunkSize` in `(arr $ chunksOf chunkSize)`. These chunks are then fed into a list `[arr [a] [b]]` of parallel arrows created by feeding chunks of the passed `ChunkSize` into the regular `parEvalN` by using `listApp`. The resulting `[[b]]` is lastly converted into `[b]` with `arr concat`.

```
1 parEvalNLazy :: (ArrowParallel arr a b conf, ArrowChoice arr, ArrowApply arr) =>
2   conf -> ChunkSize -> [arr a b] -> (arr [a] [b])
3 parEvalNLazy conf chunkSize fs =
4   arr (chunksOf chunkSize) >>>
5   listApp fchunks >>>
6   arr concat
```

```

7   where
8   fchunks = map (parEvalN conf) $ chunksOf chunkSize fs

```

6.2 parEval2

We have only talked about parallelizing the computation of arrows of the same type up until now. But sometimes we want to parallelize inhomogenous types as well. For this, we introduce a helper combinator `arrMaybe` first, that converts an arrow `arr a b` to an arrow `arr (Maybe a) (Maybe b)`.

```

1 arrMaybe :: (ArrowApply arr) => (arr a b) -> arr (Maybe a) (Maybe b)
2 arrMaybe fn = (arr $ go) >>> app
3   where
4     go Nothing = (arr $ \Nothing -> Nothing, Nothing)
5     go (Just a) = ((arr $ \ (Just x) -> (fn, x)) >>> app >>> arr Just, (Just a))

```

With this, we can now easily write `parEval2` which combines two arrows `arr a b` and `arr c d` into a new parallel arrow `arr (a, c) (b, d)`. We start by converting both arrows with `arrMaybe`, combining them with `***` into a new arrow `arr (Maybe a, Maybe c) (Maybe b, Maybe d)`. This is then replicated twice and fed into `parEvalN` to get a `arr [(Maybe a, Maybe c)] [(Maybe b, Maybe d)]`. We can then apply this arrow to the input `[(Just a, Nothing), (Nothing, Just c)]` and then extract the resulting values with `fromJust` and the `!!` operator on lists in the last step.

```

1 parEval2 :: (ArrowParallel arr a b conf,
2   ArrowParallel arr (Maybe a, Maybe c) (Maybe b, Maybe d) conf,
3   ArrowApply arr) =>
4   conf -> arr a b -> arr c d -> (arr (a, c) (b, d))
5 parEval2 conf f g =
6   (arr $ \ (a, c) -> (f.g, [(Just a, Nothing), (Nothing, Just c)])) >>>
7   app >>>
8   (arr $ \ comb -> (fromJust (fst (comb !! 0)), fromJust (snd (comb !! 1))))
9   where
10  f.g = parEvalN conf $ replicate 2 $ arrMaybe f *** arrMaybe g

```

6.3 parMap

`parMap` is probably the most common skeleton for parallel programs. We can implement it with `ArrowParallel` by repeating an arrow `arr a b` and then passing it into `parEvalN` to get an arrow `arr [a] [b]`. Just like `parEvalN`, `parMap` is 100 % strict.

```

1 parMap :: (ArrowParallel arr a b conf, ArrowApply arr) =>
2   conf -> (arr a b) -> (arr [a] [b])
3 parMap conf f =
4   (arr $ \ as -> (f, as)) >>>
5   (first $ arr repeat >>>

```

```

6   arr (parEvalN conf)) >>>
7   app

```

6.4 parMapStream

As `parMap` is 100% strict it has the same restrictions as `parEvalN` compared to `parEvalNLazy`. So it makes sense to also have a `parMapStream` which behaves like `parMap`, but uses `parEvalNLazy` instead of `parEvalN`.

```

1 parMapStream :: (ArrowParallel arr a b conf, ArrowChoice arr, ArrowApply arr) =>
2   conf -> ChunkSize -> arr a b -> arr [a] [b]
3 parMapStream conf chunkSize f =
4   (arr $ \as -> (f, as)) >>>
5   ( first $ arr repeat >>>
6     arr (parEvalNLazy conf chunkSize)) >>>
7   app

```

6.5 farm

`parMap` spawns every single computation in a new thread (at least for the instances of `ArrowParallel` we gave in this paper). This can be quite wasteful and a `farm` that equally distributes the workload over `numCores` workers (if `numCores` \leq `actualProcessorCount`, the fastest processor(s) to finish will get more tasks) seems useful.

```

1 farm :: (ArrowParallel arr a b conf, ArrowParallel arr [a] [b] conf,
2   ArrowChoice arr, ArrowApply arr) =>
3   conf -> NumCores -> arr a b -> arr [a] [b]
4 farm conf numCores f =
5   (arr $ \as -> (f, as)) >>>
6   ( first $ arr mapArr >>> arr repeat >>>
7     arr (parEvalN conf)) >>>
8   (second $ arr (unshuffle numCores)) >>>
9   app >>>
10  arr shuffle

```

The definition of `unshuffle` is

```

1 unshuffle :: Int
2   -> [a]
3   -> [[a]]
4 unshuffle n xs = [takeEach n (drop i xs) | i <- [0..n-1]]

```

, while `shuffle` is defined as:

```

1 shuffle :: [[a]]
2   -> [a]
3 shuffle = concat . transpose

```


These were taken from Eden's source code.

Do we want a farm that works on a list of arrows as well?

For this we just have to replace the mapArr with zipWithArr

6.6 farmChunk

As farm is basically just parMap with a different work distribution, it is, again, 100% strict. So we define farmChunk which uses parEvalNLazy instead of parEvalN like this:

```
1 farmChunk :: (ArrowParallel arr a b conf, ArrowParallel arr [a] [b] conf,  
2   ArrowChoice arr, ArrowApply arr) =>  
3   conf -> ChunkSize -> NumCores -> arr a b -> arr [a] [b]  
4 farmChunk conf chunkSize numCores f =  
5   (arr $ \as -> (f, as)) >>>  
6   ( first $ arr mapArr >>> arr repeat >>>  
7     arr (parEvalNLazy conf chunkSize)) >>>  
8   (second $ arr (unshuffle numCores)) >>>  
9   app >>>  
10  arr shuffle
```

7 Syntactic Sugar

To make using our new API a bit easier, we also introduce some syntactic sugar. We start with |>>>|, which is basically >>> on lists of arrows.

```
1 (|>>>|) :: (Arrow arr) => [arr a b] -> [arr b c] -> [arr a c]  
2 (|>>>|) = zipWith (>>>)
```

For the basic Arrow case, we have the *** combinator which allows us to combine two arrows arr a b and arr c d into an arrow arr (a, c) (b, d) which does both computations at once. This can easily be lifted into a parallel computation with parEval2, but requires a backend which has a implementation that does not require any configuration (hence the () as the conf parameter in the following code snippet).

```
1 (|***|) :: (ArrowParallel arr a b (),  
2   ArrowParallel arr (Maybe a, Maybe c) (Maybe b, Maybe d) (),  
3   ArrowApply arr) =>  
4   arr a b -> arr c d -> arr (a, c) (b, d)  
5 (|***|) = parEval2 ()
```

With this we can analogously to the serial &&& define the parallel |&&&|.

```
1 (|&&&|) :: (ArrowParallel arr a b (),  
2   ArrowParallel arr (Maybe a, Maybe a) (Maybe b, Maybe c) (),  
3   ArrowApply arr) =>  
4   arr a b -> arr a c -> arr a (b, c)  
5 (|&&&|) f g = (arr $ \a -> (a, a)) >>> f |***| g
```

8 Benchmarks

needed for this version of the paper? we have to find better benchmarks anyways

9 Conclusion

References

- [1] Generalising Monads to Arrows <https://jcp.org/en/jsr/detail?id=220>, November 10, 1998